



Università di Parma

Dipartimento di Ingegneria e Architettura

Intelligenza Artificiale

A.A. 2023/2024

Reinforcement Learning

Introduction

Reinforcement Learning



- ❑ Supervised (inductive) learning is the simplest and most studied type of learning
- ❑ How can an agent learn behaviors when it doesn't have a teacher to tell it how to perform?
 - The agent has a task to perform
 - It takes some actions in the world
 - At some later point, it gets feedback telling it how well it did on performing the task
 - The agent performs the same task over and over again
- ❑ This problem is called **reinforcement learning**:
 - The agent gets positive reinforcement for tasks done well
 - The agent gets negative reinforcement for tasks done poorly

Reinforcement Learning



- ❑ Learning from interaction with an environment to achieve some long-term goal that is related to the state of the environment
- ❑ The goal is defined by reward signal, which must be maximised
- ❑ Agent must be able to partially/fully sense the environment state and take actions to influence the environment state
- ❑ The state is typically described with a feature-vector

Reinforcement Learning: Exploration versus Exploitation



- ❑ We want a reinforcement learning agent to earn lots of reward
- ❑ The agent must prefer past actions that have been found to be effective at producing reward
- ❑ The agent must exploit what it already knows to obtain reward
- ❑ The agent must select untested actions to discover reward-producing actions
- ❑ The agent must explore actions to make better action selections in the future
- ❑ Trade-off between exploration and exploitation

Reinforcement Learning



◆ Learning to interact with an environment

- Robots, games, process control
- With limited human training
- Where the 'right thing' isn't obvious

◆ Supervised Learning:

- Goal: $f(x) = y$
- Data: $[< x_1, y_1 >, \dots, < x_n, y_n >]$

◆ Reinforcement Learning:

- Goal:
Maximize $\sum_{i=1}^{\infty} \text{Reward}(\text{State}_i, \text{Action}_i)$
- Data:
 $\text{Reward}_i, \text{State}_{i+1} = \text{Interact}(\text{State}_i, \text{Action}_i)$



Reinforcement Learning -Introduction

Supervised Learning:



Example

Class

Reinforcement Learning:



Situation Reward

...



Situation Reward

Playing chess: Reward comes at end of game
Ping-pong: Reward on each point scored

Reinforcement Learning - Introduction



- ❑ The goal is to get the agent to act in the world so as to maximize its rewards
- ❑ The agent has to figure out what it did that made it get the reward/punishment
 - This is known as the **credit assignment** problem
- ❑ Reinforcement learning approaches can be used to train computers to do many tasks
 - backgammon and chess playing
 - job shop scheduling
 - controlling robot limbs

Reinforcement Learning Systems



- Reinforcement learning systems have 4 main elements:
 - Policy
 - Reward signal
 - Value function
 - Optional model of the environment

Reinforcement Learning - Policy



- ❑ A policy is a mapping from the perceived states of the environment to actions to be taken when in those states
- ❑ A reinforcement learning agent uses a policy to select actions given the current environment state

Reinforcement Learning - Reward Signal



- The reward signal defines the goal
- On each time step, the environment sends a single number called the reward to the reinforcement learning agent
- The agent's objective is to maximise the total reward that it receives over the long run
- The reward signal is used to alter the policy

Value Function (1)



- ❑ The reward signal indicates what is good in the short run while the value function indicates what is good in the long run
- ❑ The value of a state is the total amount of reward an agent can expect to accumulate over the future, starting in that state
- ❑ Compute the value using the states that are likely to follow the current state and the rewards available in those states
- ❑ Future rewards may be time-discounted with a factor in the interval $[0, 1]$

Value Function (2)



- ❑ Use the values to make and evaluate decisions
- ❑ Action choices are made based on value judgements
- ❑ Prefer actions that bring about states of highest value instead of highest reward
- ❑ Rewards are given directly by the environment
- ❑ Values must continually be re-estimated from the sequence of observations that an agent makes over its lifetime

Reinforcement Learning - Formalization



- Given:
 - a state space S
 - a set of actions a_1, \dots, a_k
 - reward value at the end of each trial (may be positive or negative)
- Output:
 - a mapping from states to actions

example: Alvin (driving agent)
state: configuration of the car
learn a steering action for each state

Reactive Agent Algorithm



Repeat:

- ◆ $s \leftarrow$ sensed state
- ◆ If s is terminal then exit
- ◆ $a \leftarrow$ choose action (given s)
- ◆ Perform a

Accessible or
observable state

Policy (Reactive/Closed-Loop Strategy)



3	→	→	→	+1
2	↑		↑	-1
1	↑	←	←	←
	1	2	3	4

- A **policy** Π is a complete mapping from states to actions

Reactive Agent Algorithm



Repeat:

- ◆ $s \leftarrow$ sensed state
- ◆ If s is terminal then exit
- ◆ $a \leftarrow \Pi(s)$
- ◆ Perform a

Reinforcement Learning - Approaches



- ❑ Learn policy directly- function mapping from states to actions
- ❑ Learn utility values for states (i.e., the value function)

Value Function



- ❑ The agent knows what state it is in
- ❑ The agent has a number of actions it can perform in each state.
- ❑ Initially, it doesn't know the value of any of the states
- ❑ If the outcome of performing an action at a state is deterministic, then the agent can update the utility value $U()$ of states:
 - $U(\text{oldstate}) = \text{reward} + U(\text{newstate})$
- ❑ The agent learns the utility values of states as it works its way through the state space

Exploration



- ❑ The agent may occasionally choose to explore suboptimal moves in the hopes of finding better outcomes
 - Only by visiting all the states frequently enough can we guarantee learning the true values of all the states
- ❑ A discount factor is often introduced to prevent utility values from diverging and to promote the use of shorter (more efficient) sequences of actions to attain rewards
- ❑ The update equation using a discount factor γ is:
 - $U(\text{oldstate}) = \text{reward} + \gamma * U(\text{newstate})$
- ❑ Normally, γ is set between 0 and 1

Model-free versus Model-based



- ❑ A model of the environment allows inferences to be made about how the environment will behave
- ❑ Example: Given a state and an action to be taken while in that state, the model could predict the next state and the next reward
- ❑ Models are used for planning, which means deciding on a course of action by considering possible future situations before they are experienced
- ❑ Model-based methods use models and planning. Think of this as modelling the dynamics $p(s' | s, a)$
- ❑ Model-free methods learn exclusively from trial-and-error (i.e. no modelling of the environment)

Q-Learning



- ❑ Q-learning augments value iteration by maintaining an *estimated utility value* $Q(s,a)$ for *every action* at every state
- ❑ The utility of a state $U(s)$, or $Q(s)$, is simply the maximum Q value over all the possible actions at that state
- ❑ Learns utilities of actions (not states) \Rightarrow *model-free learning*

Q-Learning



- ❑ foreach state s
 - foreach action a
 - $Q(s,a)=0$
 - $s=\text{currentstate}$
 - do forever
 - $a = \text{select an action}$
 - do action a
 - $r = \text{reward from doing } a$
 - $t = \text{resulting state from doing } a$
 - $Q(s,a) = (1 - \alpha) Q(s,a) + \alpha (r + \gamma Q(t))$
 - $s = t$
- ❑ The *learning coefficient*, α , determines how quickly our estimates are updated
- ❑ Normally, α is set to a small positive constant less than 1

Selecting an Action



- Simply choose action with highest (current) expected utility?
- **Problem:** each action has two effects
 - yields a reward (or penalty) on current sequence
 - information is received and used in learning for future sequences
- **Trade-off:** immediate good for long-term well-being

try a shortcut - you might get lost; you might learn a new, quicker route!

Exploration policy



- ❑ Wacky approach (*exploration*): act randomly in hopes of eventually exploring entire environment
- ❑ Greedy approach (*exploitation*): act to maximize utility using current estimate
- ❑ Reasonable balance: act more wacky (exploratory) when agent has little idea of environment; more greedy when the model is close to correct

Credit Assignment Problem



- ❑ Given a sequence of states and actions, and the final sum of time-discounted future rewards, how do we infer which actions were effective at producing lots of reward and which actions were not effective?
- ❑ How do we assign credit for the observed rewards given a sequence of actions over time?
- ❑ Every reinforcement learning algorithm must address this problem

Reward Design



- ❑ We need rewards to guide the agent to achieve its goal
- ❑ Option 1: *Hand-designed reward functions*
 - This is a black art
- ❑ Option 2: *Learn rewards from demonstrations*
 - Instead of having a human expert tune a system to achieve the desired behaviour, the expert can demonstrate desired behaviour and the robot can tune itself to match the demonstration

What is Deep Reinforcement Learning?



- ❑ Deep reinforcement learning is standard reinforcement learning where a deep neural network is used to approximate either a policy or a value function
- ❑ Deep neural networks require lots of real/simulated interaction with the environment to learn
- ❑ Lots of trials/interactions is possible in simulated environments
- ❑ We can easily parallelise the trials/interaction in simulated environments
- ❑ We cannot do this with robotics (no simulations) because action execution takes time, accidents/failures are expensive and there are safety concerns



Università di Parma

Dipartimento di Ingegneria e Architettura

Intelligenza Artificiale

A.A. 2023/2024

Finite Markov Decision Processes

Slide of Karan Kathpalia

Markov Decision Process (MDP)



- ❑ Set of states S
- ❑ Set of actions A
- ❑ State transition probabilities $p(s' \mid s, a)$. This is the probability distribution over the state space given we take action a in state s
- ❑ Discount factor γ in $[0, 1]$
- ❑ Reward function $R: S \times A \rightarrow$ set of real numbers
- ❑ For simplicity, assume discrete rewards
- ❑ Finite MDP if both S and A are finite

Time Discounting



- ❑ The undiscounted formulation $\gamma = 0$ across episodes is appropriate for episodic tasks in which agent-environment interaction breaks into episodes (multiple trials to perform a task).
- ❑ Example: Playing Breakout (each run of the game is an episode)
- ❑ The discounted formulation $0 < \gamma \leq 1$ is appropriate for continuing tasks, in which the interaction continues without limit
- ❑ Example: Vacuum cleaner robot
- ❑ This presentation focuses on episodic tasks

Agent-Environment Interaction (1)



- ❑ The agent and environment interact at each of a sequence of discrete time steps $t = \{0, 1, 2, 3, \dots, T\}$ where T can be infinite
- ❑ At each time step t , the agent receives some representation of the environment state S_t in S and uses this to select an action A_t in the set $A(S_t)$ of available actions given that state
- ❑ One step later, the agent receives a numerical reward R_{t+1} and finds itself in a new state S_{t+1}

Agent-Environment Interaction (2)



- ❑ At each time step, the agent implements a stochastic policy or mapping from states to probabilities of selecting each possible action
- ❑ The policy is denoted π_t where $\pi_t(a \mid s)$ is the probability of taking action a when in state s
- ❑ A policy is a stochastic rule by which the agent selects actions as a function of states
- ❑ Reinforcement learning methods specify how the agent changes its policy using its experience

Action Selection



- At time t , the agent tries to select an action to maximise the sum G_t of discounted rewards received in the future

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1}$$

- Given current state s and action a in that state, the probability of the next state s' and the next reward r is given by:

$$p(s', r | s, a) = \Pr(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a)$$

State Transition Probabilities



- Suppose the reward function is discrete and maps from $S \times A$ to W
- The state transition probability or probability of transitioning to state s' given current state s and action a in that state is given by:

$$p(s'|s, a) = \sum_{r \in W} p(s', r|s, a)$$

Expected Rewards



- The expected reward for a given state-action pair is given by:

$$r(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = \sum_{r \in W} r \sum_{s' \in S} p(s', r | s, a)$$

State-Value Function (1)



- ❑ Value functions are defined with respect to particular policies because future rewards depend on the agent's actions
- ❑ Value functions give the expected return of a particular policy
- ❑ The value $v_{\pi}(s)$ of a state s under a policy π is the expected return when starting in state s and following the policy π from that state onwards

State-Value Function (2)



- The state-value function $v_{\pi}(s)$ for the policy π is given below. Note that the value of the terminal state (if any) is always zero.

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$

$$v_{\pi}(s) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1} | S_t = s \right]$$

Action-Value Function



- The value $q_{\pi}(s, a)$ of taking action a in state s under a policy π is defined as the expected return starting from s , taking the action a and thereafter following policy π
- q_{π} is the action-value function for policy π

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1} | S_t = s, A_t = a \right]$$

Bellman Equation (1)



- ❑ The equation expresses the relationship between the value of a state s and the values of its successor states
- ❑ The value of the next state must equal the discounted value of the expected next state, plus the reward expected along the way

Bellman Equation (2)



- The value of state s is the expected value of the sum of time-discounted rewards (starting at current state) given current state s
- This is expected value of r plus the sum of time-discounted rewards (starting at successor state) over all successor states s' and all next rewards r and over all possible actions a in current state s

$$\forall s \in S: v_{\pi}(s) = \sum_a \pi(a|s) \left[\sum_{s'} p(s', s, a) [r + \gamma v_{\pi}(s')] \right]$$

Optimality



- A policy is defined to be better than or equal to another policy if its expected return is greater than or equal to that of the other policy for all states
- There is always at least one optimal policy π_* that is better than or equal to all other policies
- All optimal policies share the same optimal state-value function v^* , which gives the maximum expected return for any state s over all possible policies
- All optimal policies share the same optimal action-value function q^* , which gives the maximum expected return for any state-action pair (s, a) over all possible policies



Università di Parma

Dipartimento di Ingegneria e Architettura

Intelligenza Artificiale

A.A. 2023/2024

Temporal-Difference Learning

What is TD learning?



- ❑ Temporal-Difference learning = TD learning
- ❑ The prediction problem is that of estimating the value function for a policy π
- ❑ The control problem is the problem of finding an optimal policy π_*
- ❑ Given some experience following a policy π , update estimate v of v_π for non-terminal states occurring in that experience
- ❑ Given current step t , TD methods wait until the next time step to update $V(S_t)$
- ❑ Learn from partial returns

Value-based Reinforcement Learning



- We want to estimate the optimal value $V^*(s)$ or action-value function $Q^*(s, a)$ using a function approximator $V(s; \theta)$ or $Q(s, a; \theta)$ with parameters θ
- This function approximator can be any parametric supervised machine learning model
- Recall that the optimal value is the maximum value achievable under any policy

Update Rule for TD(0)



- At time $t + 1$, TD methods immediately form a target $R_{t+1} + \gamma V(S_{t+1})$ and make a useful update with step size α using the observed reward R_{t+1} and the estimate $V(S_{t+1})$
- The update is the step size times the difference between the target output and the actual output

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Update Rule Intuition



- ❑ The target output is a more accurate estimate of $V(S_t)$ given the reward R_{t+1} is known
- ❑ The actual output is our current estimate of $V(S_t)$
- ❑ We simply take one step with our current value function estimate to get a more accurate estimate of $V(S_t)$ and then perform an update to move $V(S_t)$ closer towards the more accurate estimate (i.e. temporal difference)

Tabular TD(0) Algorithm



Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated
Initialize $V(s)$ arbitrarily (e.g., $V(s) = 0, \forall s \in \mathcal{S}^+$)
Repeat (for each episode):
 Initialize S
 Repeat (for each step of episode):
 $A \leftarrow$ action given by π for S
 Take action A , observe R, S'
 $V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$
 $S \leftarrow S'$
 until S is terminal

SARSA - On-policy TD Control



- ❑ SARSA = State-Action-Reward-State-Action
- ❑ Learn an action-value function instead of a state-value function
- ❑ q_π is the action-value function for policy π
- ❑ Q-values are the values $q_\pi(s, a)$ for s in S , a in A
- ❑ SARSA experiences are used to update Q-values
- ❑ Use TD methods for the prediction problem

SARSA Update Rule



- We want to estimate $q_{\pi}(s, a)$ for the current policy π , and for all states s and action a
- The update rule is similar to that for TD(0) but we transition from state-action pair to state-action pair, and learn the values of state-action pairs
- The update is performed after every transition from a non-terminal state S_t
- If S_{t+1} is terminal, then $Q(S_{t+1}, A_{t+1})$ is zero
- The update rule uses $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1})$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

SARSA Algorithm



Sarsa: An on-policy TD control algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Q-learning - Off-policy TD Control



- ❑ Similar to SARSA but off-policy updates
- ❑ The learned action-value function Q directly approximates the optimal action-value function q_* independent of the policy being followed
- ❑ In update rule, choose action a that maximises Q given S_{t+1} and use the resulting Q -value (i.e. estimated value given by optimal action-value function) plus the observed reward as the target
- ❑ This method is off-policy because we do not have a fixed policy that maps from states to actions. This is why A_{t+1} is not used in the update rule

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

One-step Q-learning Algorithm



Q-learning: An off-policy TD control algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Epsilon-greedy Policy



- At each time step, the agent selects an action
- The agent follows the greedy strategy with probability $1 - \epsilon$
- The agent selects a random action with probability ϵ
- With Q-learning, the greedy strategy is the action a that maximises Q given S_{t+1}

- ❑ The policy (i.e. the behavior learned by the agent) can be represented in a Q table (where Q stands for Q-learning, the algorithm derives from the Q function, which calculates the expected benefit of an action in the state, to create the best possible policy).
- ❑ The rows report all possible observations while the columns include all possible actions.
- ❑ The cells are therefore filled during training with values that represent the expected reward.
- ❑ The Q table only works with small observation spaces.
- ❑ When the possibilities increase, the agent must use neural networks.

Deep Q-Networks (DQN)



- ❑ Introduced deep reinforcement learning
- ❑ It is common to use a function approximator $Q(s, a; \theta)$ to approximate the action-value function in Q-learning
- ❑ Deep Q-Networks is Q-learning with a deep neural network function approximator called the Q-network
- ❑ Discrete and finite set of actions A
- ❑ Example: Breakout has 3 actions - move left, move right, no movement
- ❑ Uses epsilon-greedy policy to select actions

Q-Networks



- ❑ Core idea: We want the neural network to learn a non-linear hierarchy of features or feature representation that gives accurate Q-value estimates
- ❑ The neural network has a separate output unit for each possible action, which gives the Q-value estimate for that action given the input state
- ❑ The neural network is trained using mini-batch stochastic gradient updates and experience replay

Experience Replay



- The state is a sequence of actions and observations $s_t = x_1, a_1, x_2, \dots, a_{t-1}, x_t$
- Store the agent's experiences at each time step $e_t = (s_t, a_t, r_t, s_{t+1})$ in a dataset $D = e_1, \dots, e_n$ pooled over many episodes into a replay memory
- In practice, only store the last N experience tuples in the replay memory and sample uniformly from D when performing updates

State representation



- ❑ It is difficult to give the neural network a sequence of arbitrary length as input
- ❑ Use fixed length representation of sequence/history produced by a function $\varphi(s_t)$
- ❑ Example: The last 4 image frames in the sequence of Breakout gameplay

Q-Network Training



- ❑ Sample random mini-batch of experience tuples uniformly at random from D
- ❑ Similar to Q-learning update rule but:
 - Use mini-batch stochastic gradient updates
 - The gradient of the loss function for a given iteration with respect to the parameter θ_i is the difference between the target value and the actual value is multiplied by the gradient of the Q function approximator $Q(s, a; \theta)$ with respect to that specific parameter
- ❑ Use the gradient of the loss function to update the Q function approximator